

ML Project

UCS1603 Introduction to Machine Learning



PROJECT TEAM MEMBERS

LAKSHMI PRIYA B

185001083

RAGHAV R

185001119

NIVEDHITHA D

185001104

SAKTHI SAIRAJ

185001134



CONTENTS

01

Problem Statement

Project introduction and proposed solution

02

Dataset Details

Exploration of the dataset

03

Methodology

Step-by-step approach to solve the problem and tentative workflow

04

Model Architecture

Architecture of the model built

05

Performance Metrics

Metrics used to determine the performance of the ML Algorithm

06


Summary

Summary of the core approaches employed

Problem Statement

01

Brief discussion of the
problem statement &
proposed solution



POTHOLE & PLAIN ROAD CLASSIFICATION Using TRANSFER LEARNING & CNN

Problem Statement

Using deep learning and transfer learning techniques to solve the binary image classification problem of separating plain roads and roads with potholes by adopting an accurate model for savings in training time and computational efficiency.



Pothole



Plain road

Dataset Details

02



Exploration of the
dataset

Key Dataset Details

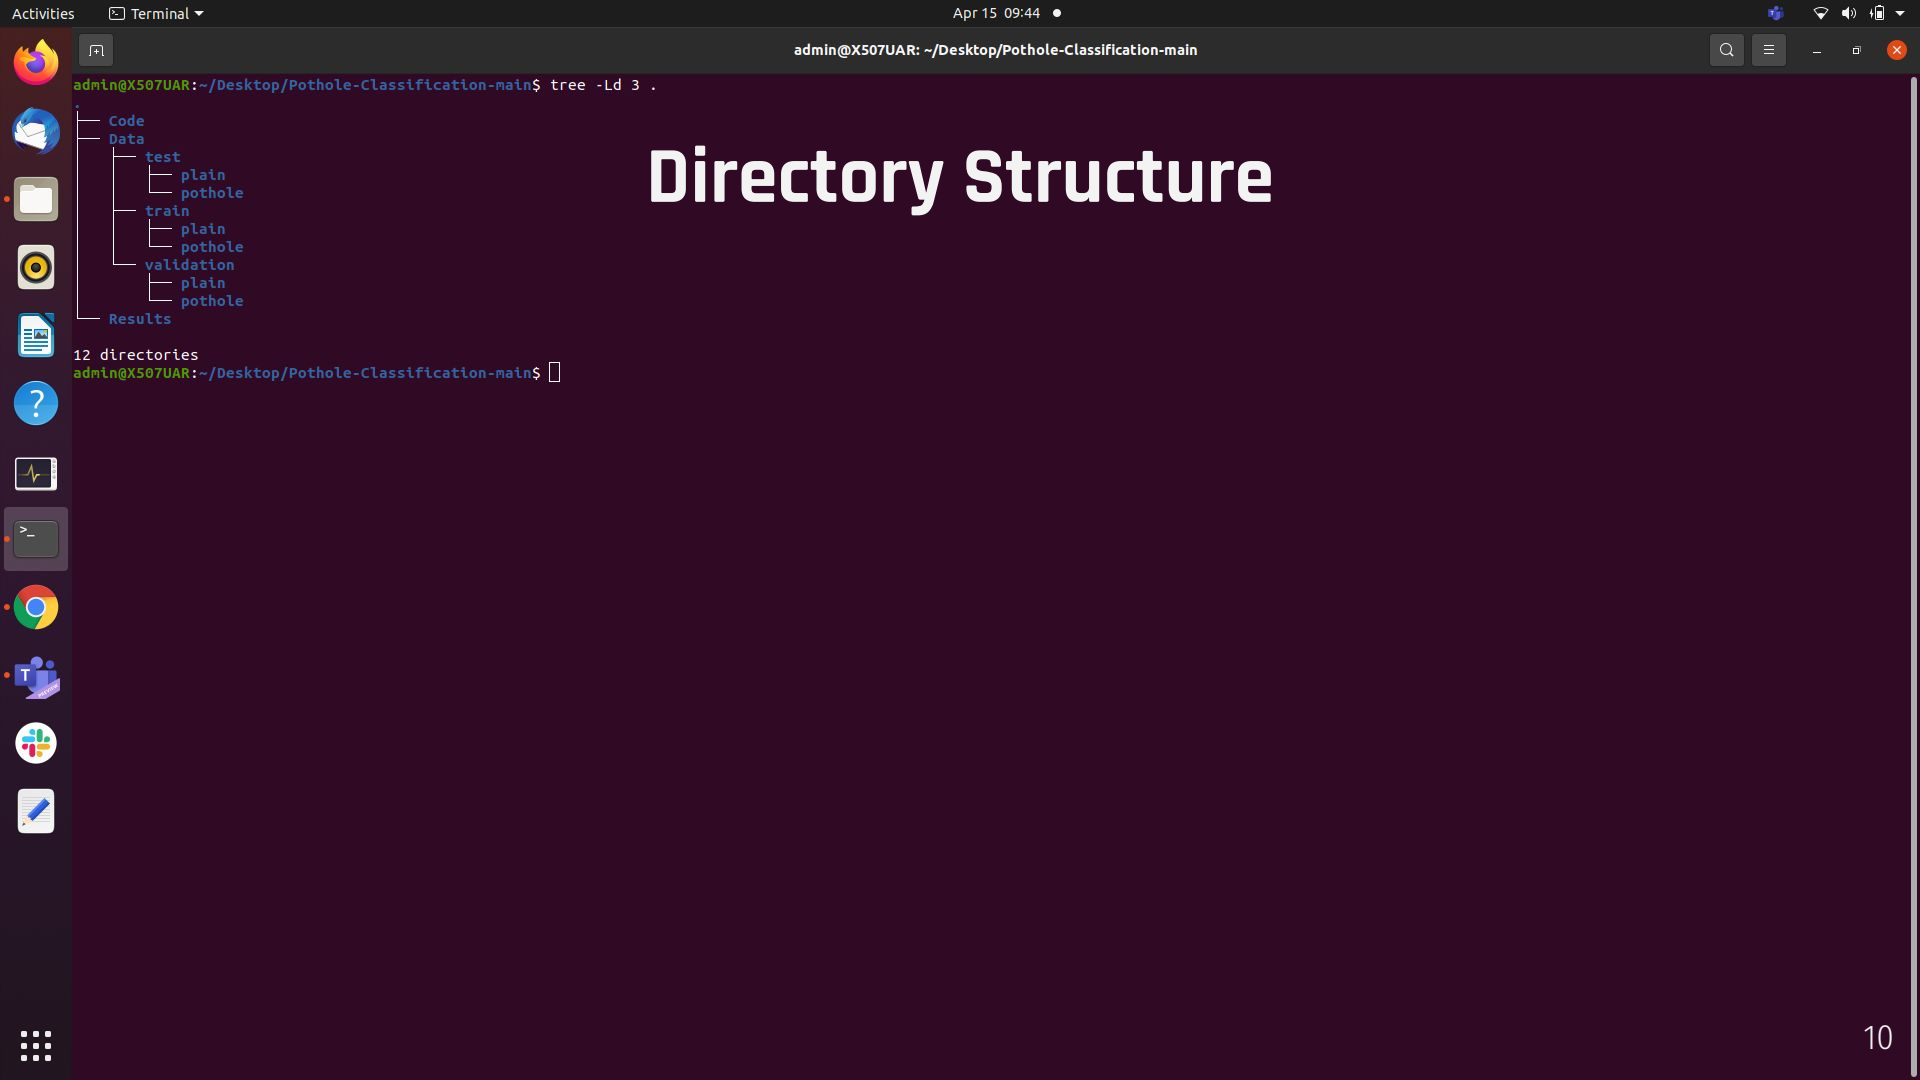
- NAME : Pothole and Plain Road Images
- SOURCE : Kaggle
- RECORD TYPE : .jpg, .jpeg , .png, .gif images
- CLASSES : Binary Classification - pothole and plain
- NUMBER OF RECORDS : 700 (350 each)
- TRAIN-VAL-TEST SPLIT : 60-20-20 ratio (210-70-70)
- FEATURES : Convolutional Base (Deep Learning)

Dataset Stats

In [10]: ▶

```
1 # Sanity checks
2 print('total training plain images :', len(os.listdir(train_plain_dir)))
3 print('total training pothole images : ', len(os.listdir(train_pothole_dir)))
4 print('total validation plain images :', len(os.listdir(validation_plain_dir)))
5 print('total validation pothole images :', len(os.listdir(validation_pothole_dir)))
6 print('total test plain images :', len(os.listdir(test_plain_dir)))
7 print('total test pothole images :', len(os.listdir(test_pothole_dir)))
```

```
total training plain images : 210
total training pothole images : 210
total validation plain images : 70
total validation pothole images : 70
total test plain images : 70
total test pothole images : 70
```



Directory Structure

```
admin@X507UAR:~/Desktop/Pothole-Classification-main$ tree -Ld 3 .
```

```
.
├── Code
├── Data
│   ├── test
│   │   ├── plain
│   │   └── pothole
│   ├── train
│   │   ├── plain
│   │   ├── pothole
│   └── validation
│       ├── plain
│       └── pothole
└── Results
```

```
12 directories
```

```
admin@X507UAR:~/Desktop/Pothole-Classification-main$
```

Methodology 03


Step-by-step approach to
solve the problem and
tentative workflow



Transfer Learning

- Transfer learning is a popular method in computer vision because it helps **build accurate models with significant savings in time**
- With transfer learning, we can solve problems from patterns that have already been learned when solving a different problem
- This way we can leverage previous learnings and avoid starting from scratch
- A pre-trained model is a model that was trained on a large benchmark dataset to solve a problem similar to the one that we want to solve
- Due to the computational cost of training such models, it is common practice to import and use models from published literature (e.g. VGG, Inception, MobileNet)

Transfer Learning




Transfer learning: idea

Instead of training a deep network from scratch for your task:

- Take a network trained on a different domain for a different **source task**
- Adapt it for your domain and your **target task**


Variations:

- Same domain, different task
- Different domain, same task



Traditional ML


- Isolated, single task learning:
 - Knowledge is not retained or accumulated. Learning is performed w.o. considering past learned knowledge in other tasks

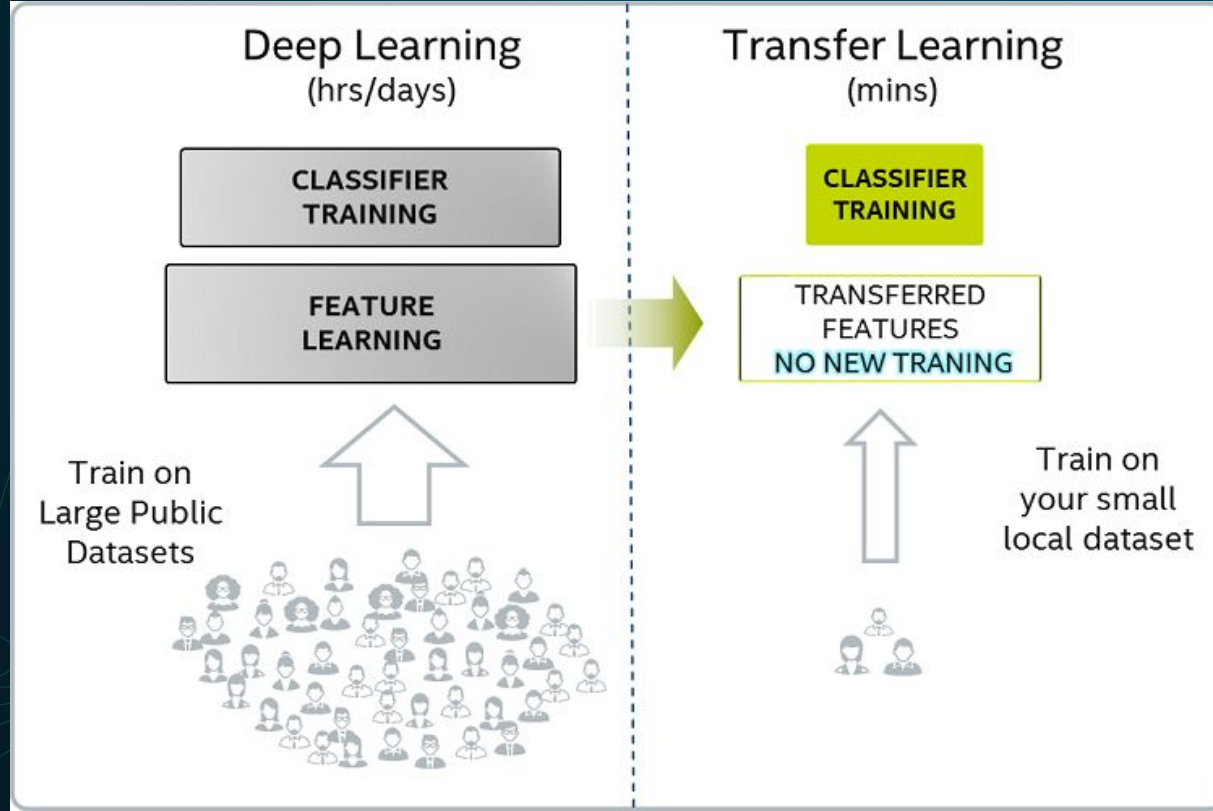


vs

Transfer Learning

- Learning of a new tasks relies on the previous learned tasks:
 - Learning process can be faster, more accurate and/or need less training data



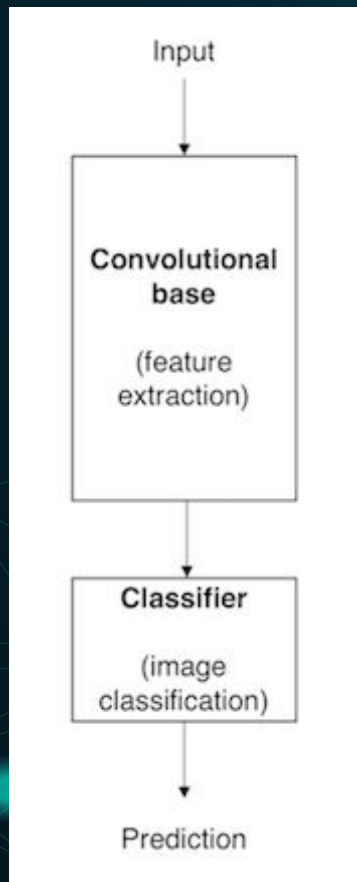




Convolutional Neural Networks

- Several pre-trained models that are used in transfer learning are based on large CNNs
- **High performance** and **ease** of training are two of the main factors driving the popularity of CNN over the last years
- A typical CNN has two parts:
 - **Convolutional base**, which is composed by a stack of convolutional and pooling layers with the main goal of generating features from the image
 - **Classifier**, which is usually composed by fully connected layers (layer whose neurons have full connections to all activation in the previous layer), has the main goal of classifying the image based on the detected features

CNN Architecture



- This deep learning model can automatically learn **hierarchical feature representations**
- This means that features computed by the first layer are general and can be reused in different problem domains, while features computed by the last layer are specific and depend on the chosen dataset and task
- *‘If first-layer features are general and last-layer features are specific, then there must be a transition from general to specific somewhere in the network’*
- As a result, the convolutional base of CNN — especially its lower layers (those who are closer to the inputs) — refer to general features, whereas the classifier part, and some of the higher layers of the convolutional base, refer to specialised features



Repurposing a pre-trained model

To repurpose a pre-trained model for our own needs, we start by removing the original classifier, then we add a new classifier that fits our purposes, and finally we have to **fine-tune our model according to one of three strategies**:

- Learning the model from scratch
- Needs a large dataset
- Lot of computational power

**Train entire
model**

- Lower layers refer to general features (problem independent), while higher layers refer to specific features (problem dependent)
- Here, we play with that dichotomy by choosing how much we want to adjust the weights of network (a frozen layer does not change during training)

**Train some layers &
leave others frozen**

- This case corresponds to an extreme situation of train/freeze trade-off
- The main idea is to keep the convolutional base in its original form and then use its outputs to feed the classifier

**Freeze
convolutional base**


Strategy 1
Train the
entire model



Strategy 2
Train some layers and
leave the others frozen



Strategy 3
Freeze the
convolutional base



Legend:

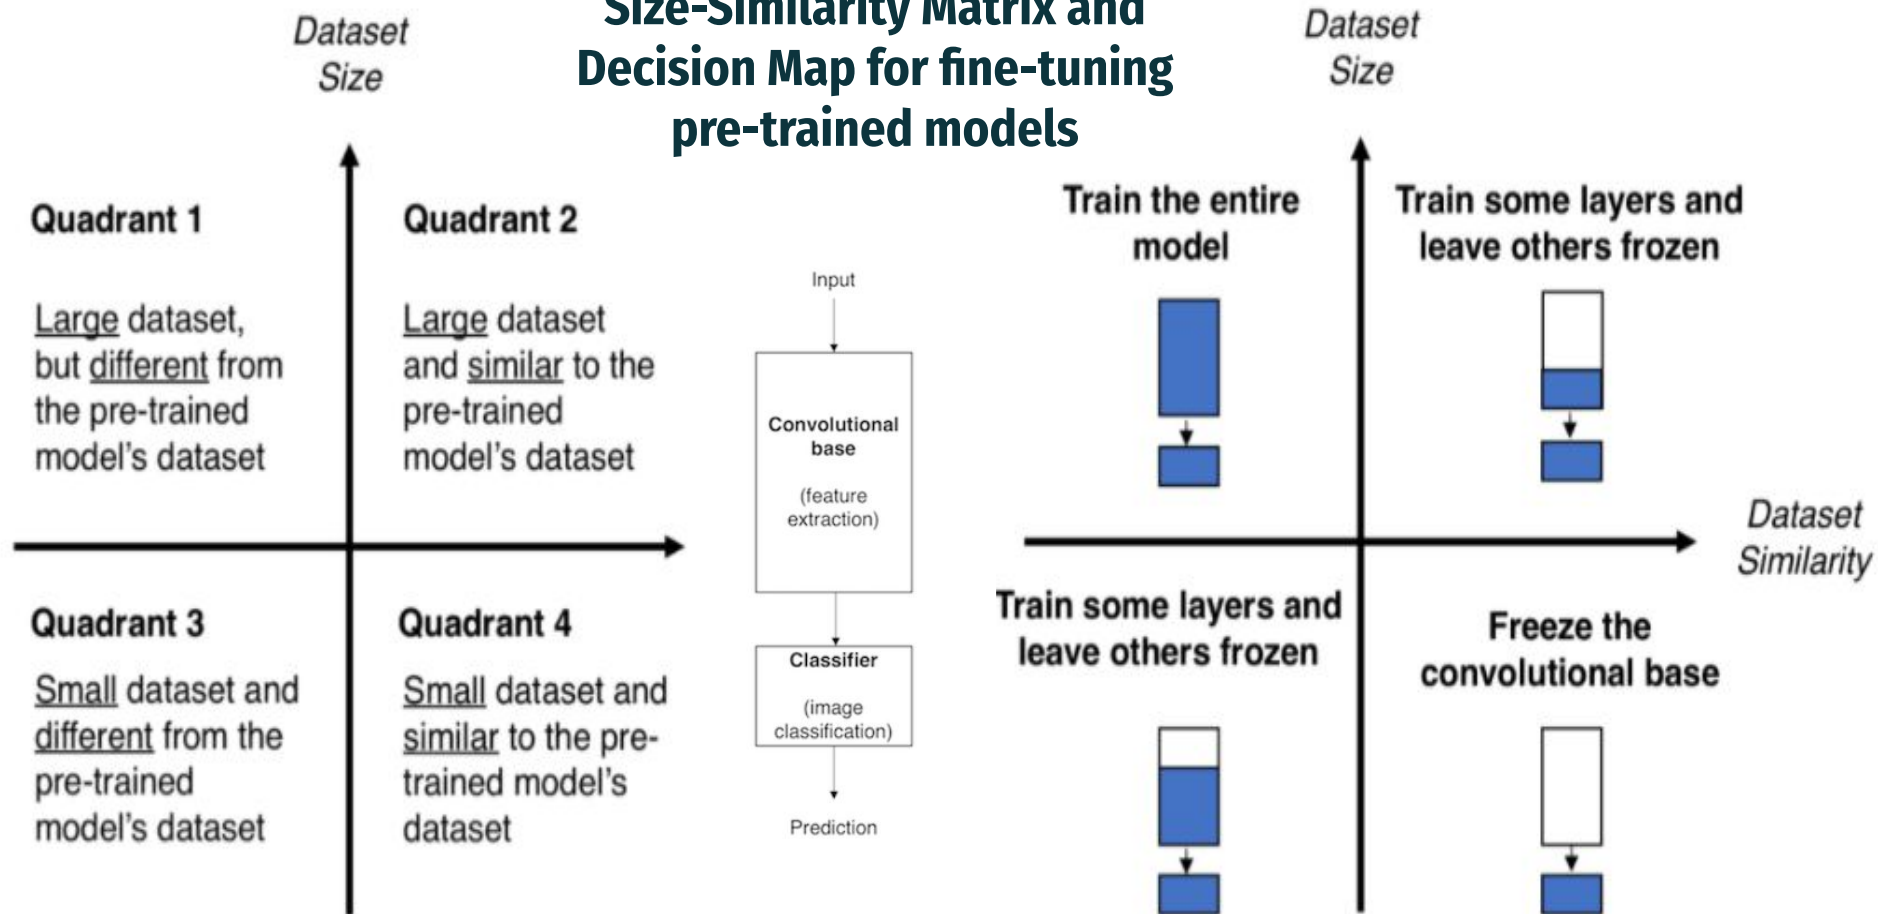




Transfer Learning Process

1. **Select a pre-trained model** based on the problem statement
2. **Classify your problem according to the Size-Similarity Matrix**
 - a. classifies your computer vision problem considering the size of your dataset and its similarity to the dataset in which your pre-trained model was trained
 - b. As a rule of thumb, consider that your dataset is small if it has less than 1000 images per class
3. **Fine-tune your model** using the Size-Similarity Matrix to guide your choice and then refer to the three options mentioned before about repurposing a pre-trained model

Size-Similarity Matrix and Decision Map for fine-tuning pre-trained models



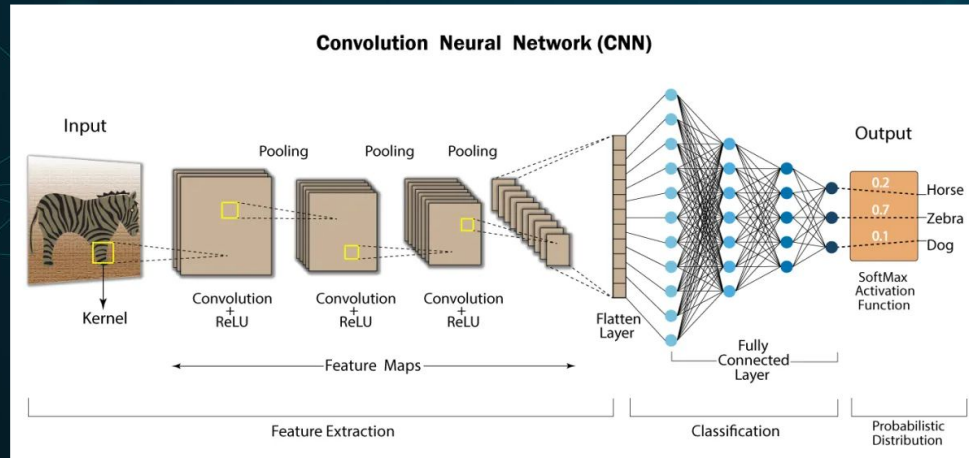
Models for image classification that result from a transfer learning approach based on pre-trained convolutional neural networks are usually composed of two parts:

1. **Convolutional base**, which performs feature extraction
2. **Classifier**, which classifies the input image based on the features extracted by the convolutional base

Classifiers on top of deep CNNs


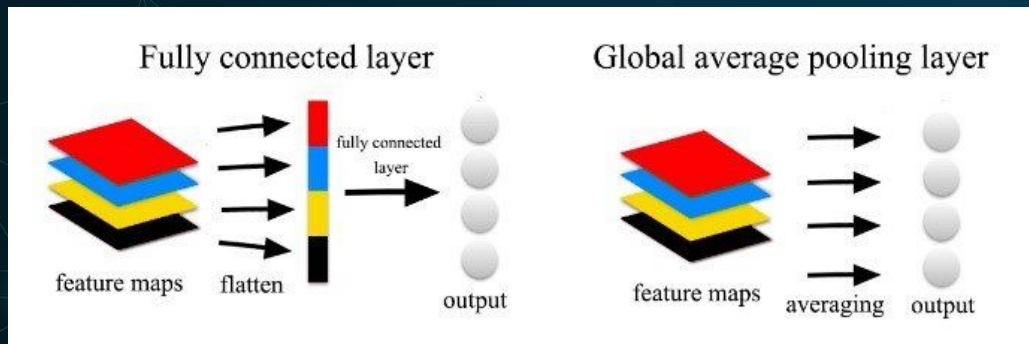
Fully-connected Layers

- For image classification problems, the standard approach is to use a stack of fully-connected layers followed by a softmax activated layer
- The **softmax layer** outputs the probability distribution over each possible class label and then we just need to classify the image according to the most probable class




Global Average Pooling

Instead of adding fully connected layers on top of the convolutional base, we add a global average pooling layer and feed its output directly into the softmax activated layer.



Linear Support Vector Machines

We can improve classification accuracy by training a linear SVM classifier on the features extracted by the convolutional base



Proposed Workflow



PREPARE DATA

Choose appropriate dataset. A smaller version of the original dataset can be used to run the models faster, which is great for people who have limited computational power



EXTRACT FEATURES

Perform feature extraction from convolutional base. These features will feed the classifiers which we want to train

Proposed Workflow



BUILDING CLASSIFIERS

1. **Fully-connected layers** - This classifier adds a stack of fully-connected layers that is fed by the features extracted from the convolutional base.
2. **Global average pooling** - The difference between this case and the previous one is that, instead of adding a stack of fully-connected layers, we will add a global average pooling layer and feed its output into a sigmoid activated layer.
3. **Linear support vector machines** - In this case, we will train a linear support vector machines (SVM) classifier on the features extracted by the convolutional base. We will use k-fold cross-validation to estimate the error of the classifier. Since k-fold cross-validation will be used, we can concatenate the train and the validation sets to enlarge our training data (we keep the test set untouched, as we did in previous cases).

Proposed Workflow



TRAIN MODEL

Train the model by fixing the epochs and batch size



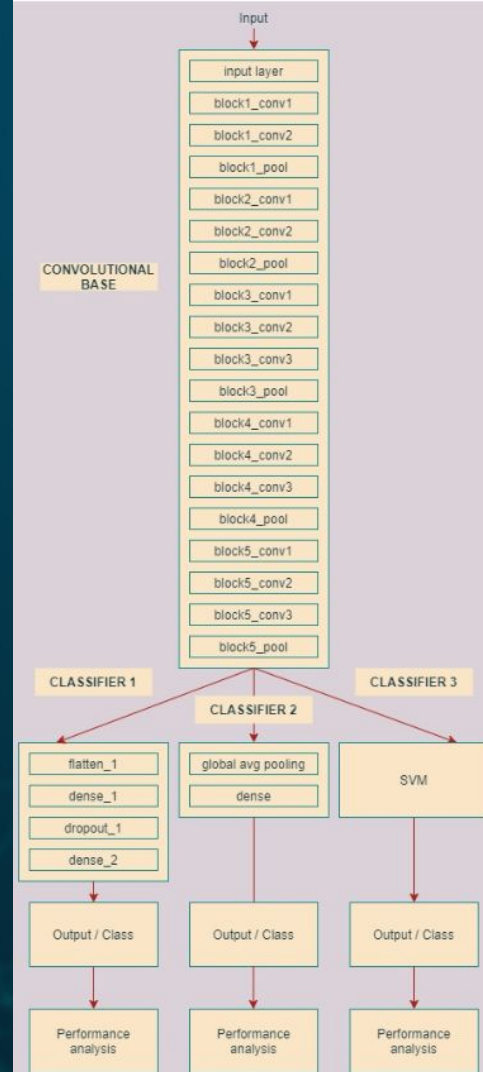

PERFORMANCE ANALYSIS

Analyse the performance of the model using various performance metrics


Model 04 Architecture

Architecture of the model
built


Architecture




Convolutional Base - VGG16 Architecture




Classifier Architecture



Fully connected layers



Global average pooling



SVM

Performance Metrics

05

Metrics used to
determine the
performance of the ML
Algorithm



Performance Metrics

“Numbers have an important story to tell. They rely on you to give them a voice.”

- Performance metrics are used to find the effectiveness of a model based on some metric using the test dataset
- Different performance metrics are used to evaluate different Machine Learning Algorithms
- The metrics that we choose to evaluate our machine learning model is very important
- Choice of metrics influences how the performance of machine learning algorithms is measured and compared

		True condition			
Total population		Condition positive	Condition negative	Prevalence = $\frac{\Sigma \text{Condition positive}}{\Sigma \text{Total population}}$	Accuracy (ACC) = $\frac{\Sigma \text{True positive} + \Sigma \text{True negative}}{\Sigma \text{Total population}}$
Predicted condition	Predicted condition positive	True positive	False positive, Type I error	Positive predictive value (PPV), Precision = $\frac{\Sigma \text{True positive}}{\Sigma \text{Predicted condition positive}}$	False discovery rate (FDR) = $\frac{\Sigma \text{False positive}}{\Sigma \text{Predicted condition positive}}$
	Predicted condition negative	False negative, Type II error	True negative	False omission rate (FOR) = $\frac{\Sigma \text{False negative}}{\Sigma \text{Predicted condition negative}}$	Negative predictive value (NPV) = $\frac{\Sigma \text{True negative}}{\Sigma \text{Predicted condition negative}}$
		True positive rate (TPR), Recall, Sensitivity, probability of detection, Power = $\frac{\Sigma \text{True positive}}{\Sigma \text{Condition positive}}$	False positive rate (FPR), Fall-out, probability of false alarm = $\frac{\Sigma \text{False positive}}{\Sigma \text{Condition negative}}$	Positive likelihood ratio (LR+) = $\frac{\text{TPR}}{\text{FPR}}$	Diagnostic odds ratio (DOR) = $\frac{\text{LR+}}{\text{LR-}}$ $F_1 \text{ score} = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$
		False negative rate (FNR), Miss rate = $\frac{\Sigma \text{False negative}}{\Sigma \text{Condition positive}}$	Specificity (SPC), Selectivity, True negative rate (TNR) = $\frac{\Sigma \text{True negative}}{\Sigma \text{Condition negative}}$	Negative likelihood ratio (LR-) = $\frac{\text{FNR}}{\text{TNR}}$	

Output: Fully-Connected Layers



Confusion Matrix:

```
[[69  1]
 [ 2 68]]
```

Accuracy: 0.9785714285714285

Specificity: 0.9857142857142858

Precision/Positive Predictive Value: 0.9855072463768116

Negative Predictive Value: 0.971830985915493

Recall/Sensitivity: 0.9714285714285714

False Positive Rate: 0.3333333333333333

False Negative Rate: 0.02857142857142857

Positive Likelihood Ratio: 2.9142857142857146


Negative Likelihood Ratio: 0.02898550724637681

Diagnostic Odds Ratio: 100.54285714285716

False Omission Rate: 0.028169014084507043

F1 Score: 0.9784172661870504



Area under ROC curve: 0.9785714285714286



Output: Performance Metrics

Fully-Connected Layers

Output: Global Average Pooling



Confusion Matrix:

```
[[67  3]
 [ 5 65]]
```

Accuracy: 0.9428571428571428

Specificity: 0.9571428571428572

Precision/Positive Predictive Value: 0.9558823529411765

Negative Predictive Value: 0.9305555555555556

Recall/Sensitivity: 0.9285714285714286

False Positive Rate: 0.375

False Negative Rate: 0.07142857142857142

Positive Likelihood Ratio: 2.4761904761904763


Negative Likelihood Ratio: 0.07462686567164178

Diagnostic Odds Ratio: 33.180952380952384

False Omission Rate: 0.06944444444444444

F1 Score: 0.9420289855072465

Area under ROC curve: 0.9428571428571428



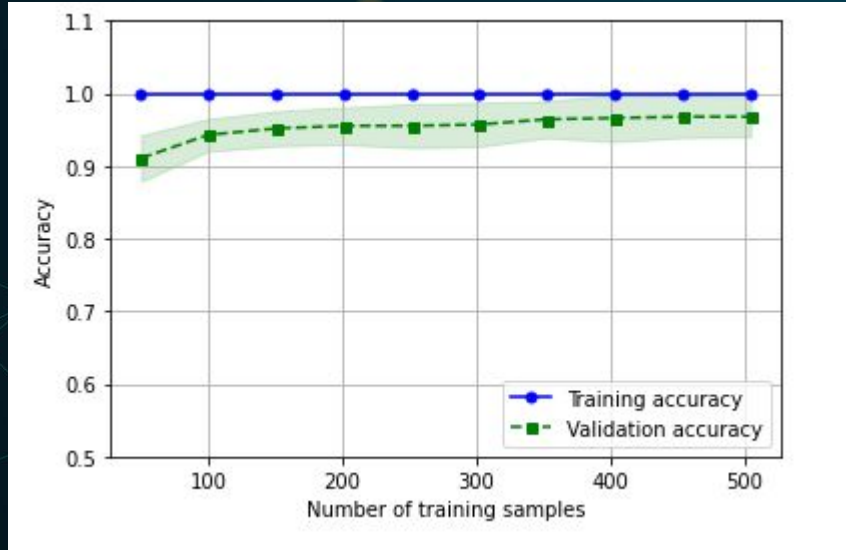
Output: Performance Metrics

Global Average Pooling

Output:

Linear Support Vector Machines

Learning Curve



Confusion Matrix:

```
[[68  2]  
 [ 3 67]]
```

Accuracy: 0.9642857142857143

Specificity: 0.9714285714285714

Precision/Positive Predictive Value: 0.9710144927536232

Negative Predictive Value: 0.9577464788732394

Recall/Sensitivity: 0.9571428571428572

False Positive Rate: 0.4

False Negative Rate: 0.04285714285714286

Positive Likelihood Ratio: 2.392857142857143


Negative Likelihood Ratio: 0.04411764705882353

Diagnostic Odds Ratio: 54.238095238095234

False Omission Rate: 0.04225352112676056

F1 Score: 0.9640287769784173

Area under ROC curve: 0.9642857142857144



Output: Performance Metrics

Linear Support Vector Machines

Comparison of the Accuracy of the Three Different Classifiers

	Fully Connected Layer	Global average Pooling	SVM
Accuracy	97.86%	94.29%	96.43%

Summary

06

Summary of the core
approaches in the
presentation

Summary

- Presented the concepts of transfer learning, convolutional neural networks, and pre-trained models.
- Defined the basic fine-tuning strategies to repurpose a pre-trained model.
- Described a structured approach to decide which fine-tuning strategy should be used, based on the size and similarity of the dataset.
- Introduced three different classifiers that can be used on top of the features extracted from the convolutional base.
- Provided a insights about each of the three classifier.
- Analysed each of the three classifiers.
- Drew ROC curve for each of the three classifiers.
- Reported the accuracy of the three classifiers.



**Thank
You**